

1

TRACK A · AI FOUNDATIONS · KIT 1 OF 20

AI Foundations & Safety: The One Hard Rule

Word-for-Word Facilitator Script: read it aloud exactly as written, or make it your own. Keyed to all 30 slides, with stage directions, a running clock, and answers to the questions staff actually ask.

Facilitator Script · 45–60 minutes

Built by Adam & Katelyn Spinozzi

A husband-and-wife team of certified educators with over 20 combined years in the classroom. Adam is a certified CTE Carpentry teacher; Katelyn is a certified K-8 teacher.

How to use this script

Everything in plain type is meant to be spoken aloud. You can read it verbatim (it's written for the ear, not the page) or glance and paraphrase once you're comfortable. Either way delivers a complete session.

Meet Ms. Rivera. Throughout the deck you'll see the same example teacher's screen: Ms. Rivera, our running composite (not a real person; she appears in every kit). When her screen is up, point the room at it; it's the model to mimic in the lab. If someone asks who she is, say exactly that: a fictional teacher we use across the series so you always see what the work looks like before you try it yourself. Her tracker chip sits top-right on every slide; the matching line under each slide cue below tells you what it says, so you always know what the room is reading.

SLIDE 0 Slide cues look like this 0:00

Advance to the slide when you reach its cue. The clock is cumulative time from the start; if your clock reads later than the cue, tighten up.

→ Stage directions look like this: *pauses, timings, what to do with your hands.*

They are never read aloud.

Boxes like this hold likely questions with suggested answers. Use them if the question comes up; skip them if it doesn't. You are never required to know more than the script.

Three delivery notes

- **Protect the lab.** The 15 hands-on minutes starting at 0:30 are the session. If you're behind, cut discussion, never the lab.
- **Let wrong AI answers happen.** If a tool fails live, point at it: that's slide 6 proving itself. Nothing you say teaches verification better.
- **Keep the tone warm and level.** No hype, no doom. You're a colleague sharing something useful and the one rule that keeps everyone safe.

Timing checkpoints

Clock	You should be...	If behind...
0:12	Starting the One Hard Rule segment (slide 10)	Skip the show-of-hands on slide 3
0:22	Starting safe-practice rewrites (slide 16)	Compress slides 12–13; the handout repeats them
0:30	Starting the lab (slide 20)	Start it anyway; run one task, not two

0:52 On First 48 Hours (slide 28)

Hold the closing lines; hand out sheets as they leave

45-minute version: run one lab task (slide 21 only), take two debrief voices instead of four, and read the shaded "45-min cut" notes where they appear.

From the founders

→ *This note is from Adam and Katelyn Spinozzi, the two certified teachers who wrote this kit. It's here to give you their perspective before you present. Sharing it is entirely your call: read parts aloud where they fit, retell them in your own words, or keep it to yourself as background. The session is complete either way.*

We're Adam and Katelyn, and here is why we believe this session matters more than any other hour of PD your staff will sit through this year: AI is already in your building, in backpacks and back pockets, whether the adults are ready or not. The teachers who will thrive with it aren't the ones who avoid it or the ones who worship it. They're the ones who understand what it actually is, follow one hard privacy rule without exception, and use the time it returns for the parts of teaching only a human can do. Get that foundation right and everything that follows in this series works. Get it wrong and no tool will fix it.

The lessons we believe this hour will leave with your staff: what these tools actually are, so nobody is intimidated by them or fooled by them; the one hard privacy rule, followed without exception, that keeps every educator safe; the habit of verifying anything factual before it's trusted; and the conviction we hold most strongly, that over-reliance is the real risk. AI can help, but only after a person can first identify the problem, understand what tools are needed to solve it, and see where AI fits in that toolbox. If the tool does that thinking for them, the muscle never develops. AI drafts. The teacher decides.

Segment 1 · Welcome & why now (slides 1–4)

SLIDE 1 Title

0:00

→ *Slide 1 is on screen as people arrive. Start on time, standing, deck behind you.*

Welcome, everyone. Today we're doing something a little different: an hour on AI: what it actually is, what it's genuinely useful for in our jobs, and the one rule every one of us needs to follow to use it safely. By the end you'll have used an AI tool with your own hands, and you'll walk out with prompts you can use tomorrow morning.

SLIDE 2 What today is (and isn't)

0:01

MS. RIVERA'S SCREEN · SO FAR Meet her: our running example teacher. Same piles on her desk as yours.

Two promises before we start. First: this is not a sales pitch. Nobody here is going to tell you AI is magic, or that you must love it. Some of the most important people in this room are the skeptics, and today will give you better reasons for your skepticism, not fewer.

Second: this is not a scare session. You may have heard AI will replace teachers. Here's our position, and it's the position of the two classroom teachers who built this training: AI is a tool. It can take over some of the robotic parts of this job: the reformatting, the first drafts, the paperwork. It cannot do the human parts: knowing your kids, building relationships, deciding what good work looks like. Today is about using the tool without ever handing it the parts of the job that make it teaching.

SLIDE 3 Why now: three numbers

0:02

MS. RIVERA'S SCREEN · SO FAR She noticed: her students and colleagues are already using these tools.

→ *Quick show of hands: "Who has used ChatGPT or a similar tool at least once, for anything?" Acknowledge the count, whatever it is. (If pressed for time, skip the hands.)*

Three numbers explain why we're spending an hour on this. In the 2023–24 school year, about one in four U.S. teachers used AI tools for their work. That's from RAND, one of the most careful research shops in education. One year later, RAND and Gallup both found it had roughly doubled: more than half of teachers were using AI. That's one of the fastest adoption curves in the history of school technology.

And the third number is about our students: Pew Research found that over half of American teens are already using AI for schoolwork. Whatever we decide in this room, AI is already in our building, in backpacks and back pockets. The only question is whether the adults are ready.

SLIDE 4 Agenda + one promise

0:03

MS. RIVERA'S SCREEN · SO FAR Her goal this hour: leave with one real task done with AI, safely.

Here's the hour: first, what these tools actually are: ten minutes, no jargon. Then the one hard rule about student privacy, the most important ten minutes of the session. Then you practice telling safe prompts from unsafe ones. Then fifteen minutes hands-on with a real tool. We close with what we'll commit to as a school, and three small things to try in the next 48 hours.

One promise from the research on professional development: PD only changes anything when you actually do the thing, and when there's follow-up. So you will type today, and this session comes with a 30-day plan; this isn't a one-and-done.

Segment 2 · What AI actually is (slides 5–9)

SLIDE 5 What a chatbot actually does

0:04

MS. RIVERA'S SCREEN · SO FAR Her mental note: it predicts words. It isn't looking anything up.

Let's demystify the thing. When you type into ChatGPT, Claude, Gemini, or Copilot, you are not searching a database of facts. These tools were built by reading enormous amounts of text, and what they learned is patterns: given these words, what words are likely to come next. That's the whole trick: extremely sophisticated prediction of likely next words.

That one fact explains almost everything you'll ever see an AI do, the impressive parts and the failures. It's why the writing sounds so fluent: fluent is exactly what it was trained to produce. And it's why it can be confidently, smoothly wrong: it isn't checking facts, it's continuing a pattern.

SLIDE 6 **Confidently wrong**

0:06

MS. RIVERA'S SCREEN · SO FAR Her habit forming: anything factual gets checked before she relies on it.

Here's how wrong it can be while sounding right. In a peer-reviewed study published last year, researchers asked a leading AI model to write literature reviews with citations. Nearly twenty percent of the citations it produced were completely made up; the papers didn't exist. And of the citations that were real, almost half contained errors. Every one of them looked perfect on the page.

This is called hallucination, and it is not a bug they're about to fix; it's a side effect of how prediction works. The professional habit that follows: **anything factual an AI gives you gets verified before you rely on it or hand it to anyone else**. Treat AI output as a draft from a source you don't fully trust, because that's what it is.

If asked "then why use it at all?" Because most of what we'd use it for isn't factual lookup. It's drafting, rewording, differentiating, formatting: tasks where you are the fact-checker by definition, because you know the content and it's your name on the work.

SLIDE 7 **The eager intern**

0:08

MS. RIVERA'S SCREEN · SO FAR Her new frame: a fast, eager intern. She stays the professional in charge.

The most useful way to hold all this in your head: AI is an eager intern. It's fast, tireless, surprisingly capable, desperate to please, and it will absolutely make things up rather than admit it doesn't know. You'd never let an intern send a parent email unread, grade without review, or meet your students unsupervised. Same intern, same rules.

And notice what that makes you. The intern doesn't replace the teacher; the intern only works because there's a professional in charge. You are the judgment in this arrangement. That's the posture for the whole series: **AI drafts, the teacher decides**.

SLIDE 8 **What it's genuinely good at**

0:09

MS. RIVERA'S SCREEN · SO FAR Her shortlist: newsletter drafts, leveled passages, rubrics, reformatting.

So what's the intern good for? The pattern-matching work: first drafts of parent newsletters and routine emails. Taking a lesson you already like and producing three reading levels of the same passage. Turning your notes into a rubric, or a word bank, or five versions of a warm-up. Reformatting: the same content as a table, a quiz, a station card. This is real time on real tasks: in the 2025 Gallup-Walton study, teachers who used AI weekly reported saving nearly six hours a week. That's the paperwork part of the job, not the teaching part.

SLIDE 9 What it can't do

0:10

MS. RIVERA'S SCREEN · SO FAR Hers alone: knowing her kids, checking facts, caring.

And here's what the intern cannot do, at any price. It has never met your students. It doesn't know that this class shuts down after lunch, that this kid just moved here, that this one will rise to a challenge and that one needs a win today. It can't verify its own facts. And it doesn't care; it produces text, not judgment, and not relationships. Everything on this slide is your job on purpose. The point of the tool is to buy back time *for* this list, never to hand this list away.

Segment 3 · The One Hard Rule (slides 10–15)

SLIDE 10 The One Hard Rule

0:12

MS. RIVERA'S SCREEN · SO FAR Her line, memorized: no student PII in a public AI tool. Ever.

→ *Slow down here. This is the most important slide in the session. Read the rule off the screen, then pause for a full three seconds.*

Everything so far was orientation. This is the part that keeps you safe.

Never put student personally identifiable information into a public AI tool. Not a name. Not a detail that could identify a child. Ever.

If you remember one sentence from this hour, it's that one. Every other choice about AI is a judgment call you're qualified to make. This one isn't a judgment call. It's the line.

SLIDE 11 What counts as student PII

0:13

MS. RIVERA'S SCREEN · SO FAR Her checklist: names, IDs, grades, IEPs, health, family details. All out.

So what counts? The obvious things: names, photos, student ID numbers, grades, addresses, family information. But also anything from a record or about a specific child: IEP contents, accommodations, discipline history, health information, test scores attached to a kid. If it's about an identifiable student, it stays out of the chat box.

SLIDE 12 The sneaky ones

0:15

MS. RIVERA'S SCREEN · SO FAR Her test before typing: could anyone who knows our school tell who this is?

Here's where well-meaning people slip. Initials plus context: "J.M. in my 3rd period with the anxiety diagnosis": in a school our size, that's identifiable. Rare details do it too: the only seventh grader who uses a wheelchair, the new student from Ukraine. You didn't type a name, but you identified a child. The test isn't "did I use a name?" The test is: **could anyone who knows our school figure out who this is?** If yes, it doesn't go in.

If asked "even in a private chat that I delete?" Yes. Deleting your view of a chat doesn't recall what the company received or logged. The rule works precisely because it has no exceptions to remember at 9pm on a Sunday.

SLIDE 13 Why the rule exists

0:17

MS. RIVERA'S SCREEN · SO FAR Her why: FERPA, plus a text box that makes oversharing effortless.

Two reasons, one legal, one practical. Legal: federal law, FERPA, protects student education records, and schools are responsible for how information from those records is shared. A public AI tool is a third party with no agreement with our district. Privacy researchers, including the Future of Privacy Forum, which publishes the vetting guidance a lot of districts use, flag exactly this: an open chat box makes it effortless to overshare, and once typed, that information has left the building. Where it's stored, how long, and whether it trains future models depends on terms of service that change without asking us.

Practical: you personally never want to be the example in the district's cautionary email. One rule, followed always, means that's never you.

Not legal advice: if a question turns on our district's specific obligations, the honest answer is "that's one for the district office." This session is educators sharing safe practice, not attorneys interpreting law.

SLIDE 14 "Public tool" vs. approved tool

0:19

MS. RIVERA'S SCREEN · SO FAR Her rule of thumb: every tool is public until the district says otherwise.

One distinction so this rule doesn't overreach. A *public* AI tool is anything you signed up for yourself: free ChatGPT, Claude, Gemini, your personal accounts. An *approved* tool is one our district has vetted and signed a data agreement for. If the district formally adopts an education AI platform with student-data protections, that changes what's allowed on that platform, and you'll be told explicitly. Until you're told: every tool is a public tool, and the rule applies.

SLIDE 15 The good news

0:21

MS. RIVERA'S SCREEN · SO FAR Her relief: the work she wants help with never needed identities anyway.

Now the good news, and it's genuinely good: the rule barely costs you anything. Almost everything AI is useful for doesn't need student identities at all. You need "a 4th grader reading below grade level," never "Marcus Chen." Which brings us to the skill that makes this easy.

Segment 4 · Safe practice (slides 16–19)

SLIDE 16 Safe swap #1: the email

0:22

Watch the same request done unsafely and safely. Unsafe: "Write an email to Jayden Miller's mom about his three missing assignments and his outburst in class Tuesday." That's a name, a family, academics, and behavior, a full student record, pasted into a stranger's text box.

Safe: "Write a warm, professional email to a parent about a middle schooler with several missing assignments and a recent difficult day in class. Firm but supportive; end by inviting a conversation." Same email. The AI never needed the child; it needed the situation. You add the name after it's back in your own hands, in your own email account.

SLIDE 17 Safe swap #2: differentiation

0:24

MS. RIVERA'S SCREEN · SO FAR Her swap: the reading level and format, never the diagnosis.

Again, with planning. Unsafe: "Rewrite this passage for Sofia R., who has a reading IEP with a 2nd-grade fluency goal." Safe: "Rewrite this passage at a 2nd-grade reading level, same key ideas, shorter sentences, and add a five-word picture-supported vocabulary list." Notice the safe version is also just a *better prompt*: the AI never needed the diagnosis, it needed the reading level and the format. Safer and more useful is the usual trade here.

SLIDE 18 De-identify like a pro: 3 moves

0:26

MS. RIVERA'S SCREEN · SO FAR Her 3 moves, automatic: strip identity, generalize, describe the need.

The skill has three moves, and you just watched all of them. One: **strip identity**: no names, no initials, no one-of-a-kind details. Two: **generalize**: "a 7th grader," "several students," "a parent." Three: **describe the need, not the child**: reading level, behavior pattern, scaffold you want. Every prompt you ever write for school passes through those three moves. They take about ten seconds, and they become automatic within a week.

SLIDE 19 You try: spot the problem

0:27

MS. RIVERA'S SCREEN · SO FAR Her eye: initials + context still identify. Rare details identify too.

→ Read each prompt off the slide; the room calls out safe or unsafe. Keep the pace brisk: 3 minutes total.
Answers: A unsafe (name + IEP). B safe. C unsafe: "the only student who just moved here from Ukraine" identifies a child without a name. D safe: aggregate, no identities.

Four prompts on the screen. For each one, call out: safe, or unsafe? And if unsafe, what's the ten-second fix?

Prompt C is the one that splits the room. If debate breaks out, that's the point landing. Resolve it with the test from slide 12: could someone who knows our school identify the child? Then the fix: "a student who recently arrived from another country and is learning English."

Segment 5 · Hands-on lab (slides 20–23)

SLIDE 20 Lab setup

0:30

MS. RIVERA'S SCREEN · SO FAR Her lab pick, one artifact all the way through: the family newsletter blurb.

→ Energy up: this is the moment the session exists for. Get devices out and pairs formed inside two minutes. Announce which tool the school is using for the lab.

Devices out. Time to actually do this. Pair up; one screen per pair is perfect. Open [TOOL NAME]. Three ground rules for the next fifteen minutes. One: no student information; we practice like we play. Two: nothing you make right now has to be good; this is a sandbox. Three: when the AI gives you something mediocre (and it will), don't settle. That's the skill we're practicing.

And keep an eye on the screen as we go: Ms. Rivera runs one artifact, her family newsletter blurb, through the whole lab, from typed prompt to reviewed result. If you ever feel lost, copy her structure.

SLIDE 21 Lab task 1: your first useful prompt

0:32

MS. RIVERA'S SCREEN · SO FAR Her newsletter blurb, v1: option A typed, draft back. Beige, sounds like nobody.

→ Give pairs 6 minutes. Circulate. If a pair is stuck, point at option A; it works for every role in the building. Deputize your power users to float. Point at Ms. Rivera's window: option A filled in with her fall book fair, and the first draft back. Don't spoil the buried detail; it pays off in task 2.

Task one. Pick whichever prompt is closest to your real job, type it, and read what comes back. Option A: "Write a warm, professional newsletter blurb for families about [any upcoming school event], under 120 words." Option B: "Create a five-question review quiz on [any topic you teach], mixed difficulty, with an answer key." Option C: "Rewrite this paragraph at three different reading levels." Then paste any paragraph from anything you already have.

And Option D comes from a real carpentry classroom, because electives and CTE deserve prompts built for their rooms too: "Create a five-question tool safety check for the miter saw station, mixed question formats, with an answer key." Swap in any station, lab, kitchen, or instrument from your own program. Go.

SLIDE 22 Lab task 2: push back

0:38

MS. RIVERA'S SCREEN · SO FAR Her newsletter blurb, v3: two pushbacks later, hers. One invented detail caught.

→ 45-min cut: skip this task; go from task 1 straight to the slide-23 debrief at 0:40. Point at Ms. Rivera's window: the same blurb, two pushbacks later, with the invented book-fair hours caught on the read-through.

Now the move that separates people who get real value from people who quit in a week: don't accept the first draft. Tell it what's wrong, like you would an intern. "Warmer." "Shorter." "A 5th grader wouldn't know three of these words — fix that." "Make question four harder and add a diagram question." Two rounds of pushback, minimum. Watch what happens.

SLIDE 23 Debrief: what did you notice?

0:44

MS. RIVERA'S SCREEN · SO FAR Her newsletter blurb's arc: beige v1, two pushbacks, usable. One error caught.

→ Take 3–4 voices, one minute each. If the room saw an AI mistake, spotlight it warmly; it's the best possible outcome of the lab.

Let's hear it. What surprised you? What was better than you expected, and what was worse? Did anyone catch it being wrong?

Two patterns worth naming. First: the first draft is rarely the good draft; the value shows up when you push back, which is why the skill is called prompting and not wishing. Second: everything you just made still needs your eyes. It's fast, it's fluent, and it's unsigned until a professional signs it. Which is the last big idea of the day.

Segment 6 · Staying human & our norms (slides 24–27)

SLIDE 24 The human stays in charge

0:45

MS. RIVERA'S SCREEN · SO FAR Her rule for what leaves her desk: nothing AI-drafted goes out unread.

Rule number one was about what goes *into* the tool. The companion habit is about what comes *out*: nothing AI-drafted reaches a student or a parent without a human reading it first. Not because it's usually wrong, but because sometimes it is, and because your name is on it. The U.S. Department of Education's guidance on AI in schools says it in three words: keep humans in the loop. In this building that means: AI drafts, the teacher decides. Always.

SLIDE 25 Honest limits: the shadow side

0:47

MS. RIVERA'S SCREEN · SO FAR Her guardrails: keep her own voice; drafts never become decisions.

Trust requires honesty, so here are the risks nobody selling AI will lead with. Over-reliance is real: if the tool writes everything, your own writing muscle atrophies, and the same goes for our students, which is why student use gets its own session. Sameness is real: AI output has a beige, everyone-sounds-alike quality, and your voice is worth protecting; keep your sentences in anything that matters. And drift is real: it's easy to slide from "AI drafts it" to "AI decides it" without noticing. The guardrails you've learned today (verify facts, review everything, keep the judgment calls) exist because these risks are genuine, not theoretical.

SLIDE 26 What about students using it?

0:49

MS. RIVERA'S SCREEN · SO FAR Her scope today: adult use only. Student use is Kit 5's hour.

The question every faculty asks. Short version for today: the major public chatbots require users to be at least 13, and minors need parent or guardian permission under the terms of service, so student use is a real policy question, not a free-for-all, and it deserves its own session; that's Kit 5 in this series. For today, our scope is the adults: our use, our privacy practice, our judgment. Get that right and we'll be ready to lead students well.

SLIDE 27 Our three commitments

0:50

MS. RIVERA'S SCREEN · SO FAR She nodded to all three. Her interim stance until policy exists.

→ *Read all three, then ask for a visible nod or thumbs-up from the room. It matters that this is agreed, not just heard.*

Here's what I'm asking us to agree to as a staff, starting today. One: no student personally identifiable information in public AI tools. Ever. Two: a human reviews anything AI-drafted before it reaches a student or family. Three: when we find something that works, we share it; a prompt that saved you an hour belongs in the staff room, not a drawer. Can I get a nod on those three?

Segment 7 · First 48 hours & close (slides 28–30)

SLIDE 28 Your first 48 hours

0:52

MS. RIVERA'S SCREEN · SO FAR Her next 48 hours: one real task, one de-ID drill, one share.

→ *Hold up the First 48 Hours sheet, then have it passed out (or point to it if pre-distributed).*

Research on PD is blunt: if you don't use this within two days, most of it evaporates. So this sheet asks for three things, each under fifteen minutes, each using something from today: a real task done with AI, a de-identification rep, and one shared discovery. That's the whole assignment. Tomorrow morning, coffee in hand, pick one.

SLIDE 29 What's next

0:53

MS. RIVERA'S SCREEN · SO FAR Her track: Kit 1 done, seven to go in the full series.

This was Kit 1 of eight in the AI Foundations track. Next is Prompting Basics: how to get consistently useful results instead of occasionally lucky ones. Schools running the full series through an AI-Ready School membership track each educator's progress, and finishing all eight earns the AI-Ready Educator Certificate of Completion, real documentation of real professional learning. And we'll keep the momentum in our regular PLC time with short follow-ups; nothing new to schedule.

SLIDE 30 Exit ticket + close

0:54

→ *Distribute exit tickets. Collect them at the door; they're the school's PD documentation, so actually collect them.*

Before you go: two minutes on the exit ticket. One thing you're taking, one worry you still have, one thing you'll try. Your worries steer what we do next, so be honest.

Last word. The tool you used today is going to keep getting better, and none of it changes why any of us took this job. AI can write a draft. It cannot greet a kid at the door by name, notice something's off, and change that kid's day. That's you. Use the tool for the paperwork — keep the teaching. Thanks, everyone.

After the room empties: collect any stray exit tickets, note the themes in the "one worry" column, and take five minutes with the 30-Day Plan; the three PLC follow-ups it contains are how this session becomes practice instead of a memory.

© Adam & Katelyn Spinozzi · AI-Ready School. Licensed to the purchasing educator (TpT single-user license); print the handouts for your own session. School-wide licensing: ai-readyschool.com. This resource is provided for informational purposes by educators, not attorneys, and is not legal advice. Consult your district's legal counsel before adopting any policy. Third-party tools named in this kit (ChatGPT, Claude, Gemini, Copilot) are referenced factually; no partnership or endorsement is implied.