

3

TRACK A · AI FOUNDATIONS · KIT 3 OF 20

AI for Planning & Differentiation

Word-for-Word Facilitator Script: read it aloud exactly as written, or make it your own. Keyed to all 30 slides, with stage directions, a running clock, and answers to the questions staff actually ask.

Facilitator Script · 45–60 minutes

Built by Adam & Katelyn Spinozzi

A husband-and-wife team of certified educators with over 20 combined years in the classroom. Adam is a certified CTE Carpentry teacher; Katelyn is a certified K-8 teacher.

How to use this script

Everything in plain type is meant to be spoken aloud. Read it verbatim (it's written for the ear) or glance and paraphrase once you're comfortable. Either way delivers a complete session. Slide cues, stage directions, and question boxes work exactly as they did in Kits 1 and 2. Run those kits first: this session assumes the two permanent rules (no student PII in public tools; a human reviews everything) and the four-part prompt formula (role, task, context, format). One prep item is non-negotiable: staff must bring one upcoming lesson. The lab is built on it.

Timing checkpoints

Clock	You should be...	If behind...
0:12	Starting the toolkit (slide 10)	Trim discussion on slides 6–8
0:24	Starting the 4-point review (slide 16)	Compress moves 3 and 4 to one minute each
0:30	Starting the lab (slide 19)	Start it anyway; the lab is the session
0:53	On First 48 Hours (slide 28)	Hold the closing lines; hand out sheets at the door

45-minute version: in the lab, differentiate one direction instead of two, take two share-out voices instead of four, and read the "45-min cut" stage notes where they appear.

Segment 1 · Welcome & the differentiation dilemma (slides 1–4)

SLIDE 1 Title

0:00

→ Slide 1 up as people arrive. Start on time.

Welcome back. Kit 1 gave us the rules of the road. Kit 2 gave us the formula: role, task, context, format. Today the formula goes to work on the hardest everyday problem in teaching: the fact that the students in your room are not all the same student. This session is about planning and differentiation, and it ends with you holding a real artifact for a real lesson you're about to teach.

SLIDE 2 The dilemma

0:01

Let's name the dilemma honestly, both halves. Half one: differentiation works. A large meta-analysis by Deunk and colleagues found small-to-moderate positive effects on achievement, and the effects were strongest when differentiation was embedded in a bigger instructional program and supported by technology. Just sorting kids into groups isn't enough; the materials have to actually differ. Half two: making the materials differ was never humanly sustainable. Three versions of every text, every practice set, every station? Nobody was failing at that out of laziness. The plan asked for more hours than the week has.

[FOUNDER STORY · Katelyn] One spoken story for this moment: a specific unit where you knew exactly which three versions your class needed and simply could not build them all. What was the lesson, what did you triage away, and which kids felt the cut? The room needs to hear that the gap was time, not intention, from someone who lived it.

SLIDE 3 Why now

0:03

Here's why this session exists now and not five years ago. As of late 2025, 61 percent of teachers were using AI in their work, up from 34 percent just two years earlier; that's from Education Week's research center. And what are they using it for? Creating lesson plans and student resources sits among the most common uses. The Gallup teacher survey adds the part that matters: among teachers who use AI for a given task, like preparing lessons or modifying materials to student levels, 60 to 84 percent say it saves them time, and weekly users save around six hours a week. The tool that attacks the time problem specifically has arrived. Today we point it at differentiation.

SLIDE 4 Agenda + one promise

0:04

The hour: first, the planning workflow, where AI fits and where it absolutely doesn't. Then the toolkit: four differentiation moves with prompts you can copy tonight. Then quality control, because a leveled text that isn't actually at level is worse than none. And then the heart of the session: you bring the lesson you were asked to bring, and you leave with one fully differentiated artifact for it. Core version, a support version, an extension, reviewed and ready to teach. That's the promise. Not examples. Your lesson.

Segment 2 · The planning workflow (slides 5–9)

SLIDE 5 Never from scratch again

0:05

Here's the stance for today, and it fits in one breath: never from scratch again. Not "never plan again." Never start from a blank page again. The blank page was always the most expensive part of planning, and it was never the part that needed your expertise. AI drafts the skeleton. You supply the judgment. Same five words as always: AI drafts, the teacher decides.

SLIDE 6 The 30-minute lesson plan

0:07

Let me show you the best study we have on exactly this. Two researchers, Powell and Courchesne, published in PLOS ONE, sat down with ChatGPT and asked it to plan a first grade science lesson. In about 30 minutes of drafting and refining, they had a standards-aligned lesson plan. Thirty minutes. That's the good news, and it's real. Now the other half, from the same study: along the way, the drafts included questionable components, missing details, and one resource that did not exist. The AI recommended it confidently. It was fabricated.

SLIDE 7 Both halves are true

0:09

Hold both halves of that study at once, because both are true. The 30 minutes is real: a working skeleton, aligned to standards, in the time it takes to eat lunch. And the fabricated resource is real too, which is why review is non-negotiable. Not review as a formality. Review as the step where the actual teaching expertise enters the document. The study's authors reached the same conclusion we teach in every kit: this workflow needs a professional in the loop, and that professional is you.

SLIDE 8 The UDL frame: generate options

0:10

One frame makes AI genuinely useful for planning instead of just fast, and it comes from Universal Design for Learning, the CAST guidelines, now in version 3.0. UDL says: design with multiple means from the start. Multiple ways to engage, multiple ways to represent the content, multiple ways for students to show what they know. A meta-analysis by Capp found UDL an effective methodology for improving learning across the board. Here's the connection: generating options was always UDL's cost problem, and generating options is the single thing AI does fastest. So the move is: ask for options in bulk, then choose like the professional who knows the room. AI generates the menu. You order for your class.

If asked "is this just learning styles again?" No, and the distinction matters. UDL doesn't sort students into types; it builds flexibility into the materials so every student has a workable path. The research base we cite is about options and access, not matching lessons to labels.

SLIDE 9 The workflow on one slide

0:11

So the planning workflow, four steps. One: draft. Use the Kit 2 formula to get the core artifact. Two: options. Ask the AI to vary it: levels, scaffolds, formats. Three: pick and adapt, because you know which option fits which kid, and the AI never will. Four: review, with a four-point checklist we'll learn before the lab. Draft, options, pick, review. That's the whole system.

Segment 3 · The differentiation toolkit: four moves (slides 10–15)
SLIDE 10 Four moves

0:12

The toolkit is four moves, and every differentiated classroom you've ever admired runs on some mix of them. Level it: the same text at different reading levels. Scaffold it: supports that let a student climb to the same target. Stretch it: depth for students who finish in four minutes. Reformat it: the same content in a different shape. If you want the research pedigree, this is Tomlinson's classic frame: differentiate by readiness, by interest, by learning profile. The first three moves serve readiness. The fourth serves interest and profile. Nothing here is new pedagogy. What's new is the price.

SLIDE 11 Move 1: level it

0:14

Move one, the workhorse: leveled texts. One passage in, three reading levels out. Here's the prompt, and notice it's just the Kit 2 formula wearing work clothes: "You are a reading specialist. Rewrite the passage below at three levels: on-grade for 6th grade, about two years below, and about two years above. Keep every key idea, the section headings, and the same core vocabulary words. At the lower level, keep sentences under 12 words and add a 5-word vocabulary box with student-friendly definitions. At the higher level, keep the length similar but raise the complexity of the reasoning, not just the vocabulary. Label each version." Then paste your passage. Ninety seconds later you have three texts that used to cost you an evening.

SLIDE 12 Check the level, don't trust the label

0:16

Now the habit that keeps move one honest: check the level, don't trust the label. The AI will happily stamp "3rd grade level" on whatever it produces. Sometimes it's right. You verify, and here's why verification is even possible: readability is genuinely measurable. Researchers led by Crossley built the CLEAR corpus, thousands of passages with reading-ease ratings grounded in teacher judgment, precisely because level is a real, testable property of a text, not a vibe. So test it, fast, three ways. Read a paragraph aloud: your ear knows your grade level. Eyeball the sentence length: if the "easier" version still runs 25-word sentences, it isn't easier. And the gold standard: try it on one actual student tomorrow and watch where they stumble. Two minutes of checking protects everything the leveling bought you.

[FOUNDER STORY · Adam or Katelyn] One spoken story for this moment: a time one of you asked AI to level a text and the label was wrong: the "easier" version that wasn't, or the vocabulary that got simpler while the sentences got longer. How did you catch it, and what's your personal quick check now? This story makes the verify habit feel like craft, not paranoia.

SLIDE 13 Move 2: scaffold it

0:18

Move two: scaffolds. Same task, same target, more handholds. The four scaffolds AI drafts well: sentence starters, word banks, worked examples, and the text for graphic organizers. The prompt: "For the writing task below, create: five sentence starters that get a struggling writer into each paragraph; a 10-word word bank with student-friendly definitions; and one fully worked example paragraph on a different topic, so students see the shape without copying the content." That last clause is the professional touch: a worked example on a different topic teaches the structure and keeps the thinking with the student.

SLIDE 14 Move 3: stretch it

0:20

Move three: extension, and let's say the quiet rule out loud: early finishers get depth, not more of the same. A second worksheet is a punishment for being fast. The prompt: "Create three extension prompts for students who have mastered the task below. Go deeper, not longer: one prompt that applies the concept to an unfamiliar context, one that asks the student to critique or find the flaw in a flawed example, and one that asks them to teach the idea to a younger student in their own words. No additional practice problems." Your strongest students stop being your most bored ones, and it cost you a minute.

SLIDE 15 Move 4: reformat it

0:22

Move four: same content, different shape. This is the UDL move: multiple means, on demand. Take the lesson content you already have and ask: "Turn the content below into three formats: a set of four station-activity directions, a deck of eight discussion cards with one question per card, and a 10-item practice set that ramps from recall to application. Same concepts in all three; label everything." One draft becomes a menu. Some days you'll use one format; some days three stations run at once. Either way, the version you never had time to make now exists.

Segment 4 · Quality control (slides 16–18)

SLIDE 16 The 4-point review

0:24

Before the lab, the checklist that stands between an AI draft and your students. Four points, every artifact, every time. One: accurate? Facts, examples, and any named resources checked; remember the fabricated resource in the PLOS ONE study. Two: at level, actually? Run the quick checks; don't trust the label. Three: fits your students? The AI wrote for a generic class; you teach a specific one. Swap the examples that won't land. Four: sounds like you? If the directions don't sound like your directions, students notice before you do. Accurate, at level, fits, sounds like you. Fail any point, fix it; that's the review.

SLIDE 17 Honest limits

0:26

The honest-limits minute, because trust lives here. AI does not know your kids. It has never met the student who shuts down at word problems or the one who reads two grades up but writes two grades down. Leveled is not the same as verified; the label is a guess until you check it. And it can invent things: activities that don't quite work, resources that don't exist. None of that is a reason to skip the tool. All of it is the reason the teacher stays in the loop. The drafting got fast. The deciding is still yours.

SLIDE 18 The equity note

0:28

One more thing before the lab, and it's bigger than this room. RAND found that in the 2023–24 school year, about a quarter of teachers were using AI for planning or teaching, and adoption lagged in higher-poverty schools. Sit with that: the schools whose students most need differentiated materials are the least likely to be getting AI's help making them. These tools should narrow gaps, not widen them. In this building, the way we make that true is simple: everyone learns the moves, everyone shares the prompts, and the differentiation reaches every roster, not just the early adopters'.

[YOUR TAKE · Adam & Katelyn] A plainly stated shared position for this moment: what does "AI as an equity tool" mean to you in practice, and where's your line? For example: which students in your own rooms got more of your differentiated materials once drafting got cheap, and what do you refuse to let AI decide about a struggling learner? Staff trust the equity point more when it's personal, not abstract.

Segment 5 · The lab: differentiate your real lesson (slides 19–24)

SLIDE 19 Lab setup

0:30

→ *Devices out, pairs formed inside two minutes. Announce the tool. Anyone without a lesson borrows a partner's or uses next week's plan book. Energy up; this is the session.*

Lab time, and today the material is yours: the upcoming lesson you brought. By the end of three steps you'll have a core artifact, a support version, and an extension, reviewed and teachable this week. Ground rules ride along: no student information in the prompt, ever. Describe needs, not names: "students reading below grade level," never a child. And nothing goes to a student until the 4-point review says so.

[LOCAL EXAMPLE · Adam or Katelyn] Drop in one real differentiation prompt you actually use, exactly as you run it (Adam: something from the shop, like leveling a tool-safety procedure into a plain-language version and a full-spec version for advanced students; Katelyn: something from the K-8 room, like turning one read-aloud into three center activities). Read the prompt itself out loud. One real prompt from a real classroom tells the room this isn't theoretical.

SLIDE 20 Lab step 1: draft the core

0:33

→ *5 minutes. Circulate. Weak drafts usually lack context; nudge with "what would a substitute need to know about this lesson?"*

Step one: draft the core artifact with the Kit 2 formula. Pick the one artifact your lesson most needs: the reading, the practice set, the directions, the exit check. Role, task, context, format, and be generous with context: grade, topic, what students just learned, what they always trip on. Run it, read it like an editor, and get it to eighty percent. Don't polish; the next step is where today's session earns its name.

SLIDE 21 Lab step 2: two directions

0:38

→ *7 minutes. 45-min cut: one direction instead of two; let each teacher pick support or extension based on their roster.*

Step two: differentiate it, two directions. First a support move: level it down or scaffold it, your choice, using the toolkit prompts on your handout. Then an extension move: stretch it with depth prompts. Two follow-up prompts in the same chat; the AI already has your artifact. This is the moment to notice the arithmetic: you are building the versions that used to cost an evening, inside seven minutes, for a lesson you're actually about to teach.

SLIDE 22 Lab step 3: the 4-point review

0:45

→ 4 minutes. Insist on this step; skipping review is the failure mode of the whole kit. Ask pairs to swap and check each other's level labels.

Step three: run the 4-point review on all three versions. Accurate? At level, actually? Read a few sentences of the support version aloud to your partner; their ear is your first check. Fits your students? Sounds like you? Fix what fails, right now, while the chat is open. An artifact that passes all four is done. Not AI-done. Done done.

SLIDE 23 Share-out

0:48

→ 3–4 voices, one minute each. 45-min cut: two voices. Prioritize anyone whose review caught a real problem.

Let's hear a few. What did you build, which direction surprised you, and did the review catch anything? If the AI got a level wrong or invented something, tell that story first; it's the most useful minute in the room.

SLIDE 24 What you just built

0:50

Look at what just happened. Nineteen minutes ago you had a lesson plan and good intentions. Now you have a core artifact, a support version, and an extension, all reviewed by the one person qualified to review them. Remember the meta-analysis from the top of the hour: differentiation works best embedded and technology-supported. That's not a description of some other school anymore. That's what's on your screen.

Segment 6 · Making it stick (slides 25–27)

SLIDE 25 One artifact a week

0:51

Here's the habit that makes today compound: one differentiated artifact per week. Not a differentiated everything; that's how burnout dressed up as ambition. One lesson a week gets the full treatment: core, support, extension, review. In a semester that's fifteen lessons deep with materials for every kind of learner on your roster, built in the margins of time you already have. Small, steady, sustainable. That's the whole strategy.

SLIDE 26 Commitment #5

0:52

Our running list grows by one. Kit 1 gave us three commitments; Kit 2 added the staff prompt doc. Today, commitment five: nothing AI-drafted reaches a student until it passes the 4-point review. Accurate, at level, fits, sounds like you. And when a differentiation prompt works, it goes in the staff doc with the others, so the fourth grade team isn't rebuilding what the third grade team perfected. Can I get a nod on that?

SLIDE 27 What's next

0:53

Next in the series: Kit 4, AI for Assessment. Today you built the materials for the range of learners in your room; next time we build the ways to find out what they learned: question banks, rubrics, feedback that doesn't eat your Sunday. Your differentiated artifact from today comes with you; it's about to need an exit ticket.

Segment 7 · First 48 hours & close (slides 28–30)

SLIDE 28 Your first 48 hours

0:54

Three small things within two days, each under fifteen minutes. Teach with what you built today, and jot one line about what the support kids and the stretch kids actually did. Level one more text for next week and run the quick level check on it. And post your best differentiation prompt to the staff doc. Tomorrow morning, coffee in hand, pick one.

SLIDE 29 Exit ticket

0:55

→ *Distribute exit tickets; collect at the door. They're the school's PD documentation.*

Two minutes on the exit ticket: the move you'll use first, the artifact you built, and the differentiation task you still want help with. Your answers steer the follow-ups.

SLIDE 30 Close

0:56

Last word. Differentiation was never a good-intentions problem. It was a time problem. As of today, the time problem just got smaller. The knowing-your-kids part was always yours, and no tool touches it. Go teach the lesson you just planned. Thanks, everyone.

© Adam & Katelyn Spinozzi · AI-Ready School. Licensed for use within the purchasing school. This resource is provided for informational purposes by educators, not attorneys, and is not legal advice. Consult your district's legal counsel before adopting any policy. Third-party tools named in this kit (ChatGPT, Claude, Gemini, Copilot) are referenced factually; no partnership or endorsement is implied.