

4

TRACK A · AI FOUNDATIONS · KIT 4 OF 20

AI for Assessment: Rubrics, Feedback, and Question Banks

Word-for-Word Facilitator Script: read it aloud exactly as written, or make it your own. Keyed to all 30 slides, with stage directions, a running clock, and answers to the questions staff actually ask.

Facilitator Script · 45–60 minutes

Built by Adam & Katelyn Spinozzi

A husband-and-wife team of certified educators with over 20 combined years in the classroom. Adam is a certified CTE Carpentry teacher; Katelyn is a certified K-8 teacher.

How to use this script

Everything in plain type is meant to be spoken aloud. Read it verbatim (it's written for the ear) or glance and paraphrase once you're comfortable. Either way delivers a complete session. Slide cues, stage directions, and question boxes work exactly as they did in Kits 1 and 2. Run Kits 1 and 2 before this one: today leans hard on Kit 1's privacy rules and Kit 2's prompt formula and iteration rounds.

Meet Ms. Rivera. Throughout the deck you'll see the same example teacher's screen: Ms. Rivera, our running composite (not a real person; she appears in every kit). When her screen is up, point the room at it; it's the model to mimic in the lab. If someone asks who she is, say exactly that: a fictional teacher we use across the series so you always see what the work looks like before you try it yourself. Her tracker chip sits top-right on every slide; the matching line under each slide cue below tells you what it says, so you always know what the room is reading.

THE STANCE OF THIS WHOLE SESSION

AI drafts assessment materials. It never grades students, never sees student work with names attached, and never makes decisions about students. If you say only one thing three times today, say this.

Timing checkpoints

Clock	You should be...	If behind...
0:12	Starting the workshops (slide 9)	Trim discussion on slides 5–7
0:24	On the red lines (slide 15)	Never trim slides 15–16; trim slide 13 instead
0:30	Starting the lab (slide 18)	Start it anyway; the lab is the session
0:54	On First 48 Hours (slide 28)	Hold the closing lines; hand out sheets at the door

45-minute version: build a 6-question bank instead of 10, run one rubric iteration round instead of three, take two share-out voices instead of four, and read the "45-min cut" stage notes where they appear.

From the founders

→ This note is from Adam and Katelyn Spinozzi, the two certified teachers who wrote this kit. It's here to give you their perspective before you present. Sharing it is entirely your call: read parts aloud where they fit, retell them in your own words, or keep it to yourself as background. The session is complete either way.

We're Adam and Katelyn, and here is why this session matters for what happens in your classrooms: feedback is among the highest-leverage things a teacher does, and grading is the pile that crowds it out. Those hours land on evenings and weekends the contract never sees, and they crowd out the one thing the research says moves students most: timely, specific, personal feedback. AI can shrink that pile, but only on one condition, and the condition is this whole session: AI drafts assessment materials. It never grades students, never sees student work with names attached, and never makes decisions about students.

The lessons we believe this hour will leave with your staff: how to turn your own notes about good work into rubrics students can actually read; how to build question banks and feedback starters that sound like a teacher wrote them; the red lines that keep AI out of grading entirely; and the habit that makes all of it safe, that nothing reaches a student until you've personalized it and signed it. The judgment calls stay yours. That isn't a limitation of the tool. It's the design.

Segment 1 · Welcome & the grading mountain (slides 1–4)

SLIDE 1 Title

0:00

→ Slide 1 up as people arrive. Start on time.

Welcome back. So far in this series we've set the rules of the road, learned to write prompts that actually work, and put AI to work on planning. Today we walk toward the biggest pile on your desk: assessment. Rubrics, question banks, and feedback language. And before we do anything else, one bright line, stated up front: AI never grades your students. Not today, not in the lab, not ever in this building. It drafts materials. You make every call about every kid. Hold that; we'll come back to it.

SLIDE 2 The grading mountain

0:01

MS. RIVERA'S SCREEN · SO FAR Her pile: 5 hours a week of grading and feedback, most after dinner.

Two numbers first, both from the Merrimack College and EdWeek Research Center survey of how teachers actually spend their time. Teachers work about 54 hours a week, and less than half of that is spent directly teaching. Inside that week sits the pile: the median teacher spends 5 hours a week on grading and giving feedback. Five hours, every week, most of it on evenings and weekends the contract never sees. That's the mountain this session goes after. Not by grading for you. By making the materials around grading fast to build.

→ If you have your own version of this pile (most presenters do), sixty seconds on it lands well here: what the stack was, how long it took, what was left of the evening when it was done. The room trusts the session more when the presenter has carried the pile too.

SLIDE 3 The paradox

0:03

MS. RIVERA'S SCREEN · SO FAR Her gap: feedback is her highest-leverage move, and AI never touched it.

Now the strange part. Feedback is one of the most powerful things we know how to do in education. Hattie and Timperley's famous review put its average effect size at 0.79, which is roughly double the typical educational intervention. The same review is equally clear that the power swings widely depending on how the feedback is given; feedback done poorly can even do harm. So it's high-leverage and skill-dependent. And here's the gap: RAND found that in the 2023–24 school year, of the teachers who used AI at all, only about fifteen percent or fewer used it for writing assessments. It is the least-tapped use of these tools, sitting directly under the highest-leverage work we do. That's why this hour exists.

SLIDE 4 Agenda + one promise

0:04

MS. RIVERA'S SCREEN · SO FAR Her goal: a rubric and question bank for a real upcoming assessment.

The hour: first, a few minutes on what actually makes assessment materials work, because the research here is old, deep, and useful. Then three short workshops: rubrics, question banks, and feedback language, each with a worked prompt you can steal. Then the red lines, in writing, so nobody ever has to guess. And then the lab, the longest one yet: you pick a real assessment coming up on your actual calendar, and you build for it. That's the promise. You leave with a rubric and a question bank for something you're really about to teach. Not samples. Inventory.

Segment 2 · What makes assessment materials work (slides 5–8)

SLIDE 5 Formative assessment works

0:05

MS. RIVERA'S SCREEN · SO FAR Her formative habit: check understanding, adapt. AI drafts the materials.

Three pieces of evidence, quickly, because they explain what we build in the lab. First: checking understanding as you go, and adapting, raises achievement. That's Black and William's "Inside the Black Box," resting on roughly 250 studies. It's one of the most solid findings we have. But notice the catch for your workload: assessing day by day means you need more checks, more questions, more often. Formative assessment is partly a materials problem. And drafting materials is exactly the thing AI is fast at.

SLIDE 6 Rubrics work

0:07

MS. RIVERA'S SCREEN · SO FAR Her rubric bar: students can read it. That's what makes the research pay.

Second: rubrics aren't just paperwork for the binder. A 2023 meta-analysis by Panadero and colleagues found that using rubrics has a positive, moderate effect on student performance, on self-regulated learning, and on self-efficacy. When students can see the target, they aim at it, and they start coaching themselves toward it. A good rubric quietly does some of your teaching for you. The catch, again: writing one that students can actually read takes real time. Materials problem.

SLIDE 7 Quality beats quantity

0:09

MS. RIVERA'S SCREEN · SO FAR Her question bar: quality beats quantity; the misconception does the work.

Third, and this one is professional judgment more than a citation: question quality matters more than question quantity. A sharp question does diagnostic work. A multiple-choice distractor built from a real misconception tells you exactly which students hold that misconception. Five questions like that beat twenty filler questions from a workbook. The reason most of us settle for filler isn't taste; it's time. Writing sharp questions is slow. Drafting speed is what makes quality affordable.

SLIDE 8 The division of labor

0:10

MS. RIVERA'S SCREEN · SO FAR Her division of labor: AI drafts the materials, she makes every judgment.

So here's today's division of labor. AI is fast at pattern work: rubric descriptors, question stems, distractors, feedback phrasing. You are unreplaceable at judgment: what proficient actually means in your classroom, whether a question is fair to the kids in front of you, and what this particular student needs to hear next. Every minute AI saves you on drafting is a minute that can move to judgment. Same five words as always: AI drafts, the teacher decides. In assessment, those words carry the most weight they'll ever carry.

Segment 3 · Three workshops (slides 9–14)

SLIDE 9 Workshop #1: the rubric

0:12

Workshop one: rubrics. Before: "Make a rubric for my essay assignment." You can guess what that gets: four columns of fog. After: "You are an experienced 7th grade ELA teacher. Turn my notes into a 4-level analytic rubric, with levels named Beginning, Developing, Proficient, and Advanced, for a persuasive essay. Criteria: claim, evidence, organization, conventions. Write every descriptor in language a 7th grader can read, and make each level describe what the work does, not what it lacks. Format: a table with criteria down the side. Then ask me up to three questions about anything unclear in my notes." And then you paste your notes: the messy ones, the half-sentences about what an A looks like. Your notes are the context. That's Kit 2's formula doing assessment work.

SLIDE 10 The hand-it-to-a-student test

0:14

MS. RIVERA'S SCREEN · SO FAR Her rubric test: hand it to a student. 'Limited understanding' fails.

Here's the quality bar for every rubric you'll ever draft this way, and it's one sentence: could you hand it to a student? Read one descriptor aloud. If a student reading it could tell you exactly what to do to move up one level, it passes. If it says "demonstrates limited understanding," it fails, because no ten-year-old ever revised an essay off the word "limited." When a descriptor fails, that's an iteration round, not a defeat: "Rewrite level two so a student knows exactly what moving to level three requires." The rubric research from a minute ago only pays off if students can read the thing. That test is how you make sure they can.

SLIDE 11 Workshop #2: the question bank

0:16

MS. RIVERA'S SCREEN · SO FAR Her bank prompt: topic plus the confusion her students actually have.

Workshop two: question banks. Some of you will recognize the seed of this prompt from Kit 2's diagnosis drill; today it grows up. Before: "Write ten quiz questions about photosynthesis." After: "You are a middle school science teacher. Build a 10-question bank on photosynthesis for 7th graders who commonly confuse it with respiration. Mix: six multiple choice, two short answer, two extended response. In each multiple-choice question, make one distractor the respiration confusion. Include a full answer key, and after it, explain in one line why each distractor is wrong. Label each question with the skill it checks." The special ingredient is the misconception. You know what your students mix up; the model doesn't until you say so. Feed it the confusion and the distractors start doing diagnostic work.

SLIDE 12 Vet everything

0:18

MS. RIVERA'S SCREEN · SO FAR Her vetting stance: 97% on-topic still isn't 100%. Every key gets checked.

How good are the drafts? Honest answer: good, and not good enough to skip you. In a 2023 study, Elkins and colleagues had teachers rate questions generated by large language models: 97 percent were judged on-topic, and usefulness averaged 3.6 out of 4. Those numbers are why this workflow is worth your time. But read the fine print with me: that same study treats teacher vetting as integral, not optional. Some questions were still wrong, ambiguous, or off-level. So the deal is fixed forever: every bank gets a vetting pass, every answer key gets checked, and in today's lab you will each catch and cut at least one weak question by hand. On purpose. That rep is the habit.

SLIDE 13 Workshop #3: the honest evidence on feedback

0:20

MS. RIVERA'S SCREEN · SO FAR Her read of the evidence: AI feedback is a floor. She is the ceiling.

Workshop three: feedback, and I'll give you the evidence straight, both halves. Half one: Steiss and colleagues, 2024, compared feedback on student writing from well-trained humans and from ChatGPT. The trained humans won on nearly every dimension. Your feedback, done well, is better than the machine's. Half two: Kaliisa and colleagues, 2025, pooled 41 studies with over four thousand students and found no statistically significant difference in learning outcomes between AI-generated and human feedback. Both findings are real. Here's our read: AI feedback language is a serviceable floor and a very fast draft. You are the ceiling. So we don't hand feedback over to the machine, and we don't pretend drafting help is worthless. We use it where it's strong: the common patterns.

If asked "doesn't the no-difference finding mean AI could just do the feedback?" Two answers. First, the meta-analysis measured pooled learning outcomes, not trust, motivation, or relationships, and feedback is a relationship act. Second, our red line doesn't come from the effect sizes; it comes from who is responsible for students. That's you, and it isn't delegable.

SLIDE 14 Feedback starters: the workflow

0:22

MS. RIVERA'S SCREEN · SO FAR Her starters: drafted for patterns, personalized before any student sees one.

So here's the workflow, and notice it never needs student work at all. You already know your most common patterns; you write the same three comments forty times every stack. Describe the pattern, not the student: "Here are the three most common patterns in my students' persuasive essays: claims with no evidence, evidence with no explanation, and conclusions that introduce brand-new arguments. For each, draft a 2-to-3 sentence feedback comment that names something the writer did well, names the pattern plainly, and gives one concrete next step. Warm, specific, readable by a 7th grader. I'll personalize every comment before a student sees it." That last line is the contract. AI drafts the starter. You add the sentence only you can write, the one that proves you read their work, and you deliver it. Nothing auto-graded. Nothing unreviewed. Ever.

Segment 4 · The red lines & honest limits (slides 15–17)

SLIDE 15 The red lines

0:24

MS. RIVERA'S SCREEN · SO FAR Her red lines: AI never grades, never sees names, never decides.

→ *Dark slide. Slow down. This is the one slide nobody may leave without.*

Three red lines, in writing, so nobody guesses. One: AI never grades. No score, no mark, no decision about a student ever comes from a machine in this building. Two: no student names, and no identifying student work, in public AI tools. Ever. That's Kit 1's hard rule wearing its assessment clothes. Three: feedback that reaches a student always passes through you. Every comment gets your eyes and your edits before it gets their eyes. These aren't our opinions about the technology. They're our promises about our students. If a parent asks "is AI grading my kid?", every person in this room can now say no, and say what we actually do instead.

SLIDE 16 The anonymous-excerpt rule

0:26

MS. RIVERA'S SCREEN · SO FAR Her excerpt rule: anonymous or described, never pasted with identity.

One careful clarification, because it's the question you're already forming: can I ever paste a piece of student work in, to get feedback drafting help? Answer: only an anonymous excerpt, and only if it passes Kit 1's test: zero identifying details. No name, no classmates' names, no teacher names, and no story details a neighbor could recognize: not the injury, not the family trip, not the thing that happened at recess. If you're even slightly unsure, don't paste. Describe the pattern instead: "students keep confusing X with Y," "this writer repeats the claim instead of explaining the evidence." Pattern-description is always safe, and most of the time it works just as well, because the comment you need is about the pattern anyway. When in doubt, describe, don't paste.

SLIDE 17 Honest limits

0:28

MS. RIVERA'S SCREEN · SO FAR Her red lines, adopted out loud, with the presenter's own.

Three honest limits before we build. First, the big one: AI cannot know what a student was trying to do. It sees the words on the page. You saw the six weeks: the false starts, the breakthrough on Tuesday, the kid who finally attempted a counterargument. Half of good feedback speaks to the attempt, and the machine can't see attempts. Second: rubric language drifts generic fast. "Demonstrates strong understanding" could describe any assignment since the invention of school. The fix is always your assignment's specifics: name the text, the skill, the moves you taught. Third: question banks contain errors, full stop. In the 2025 Gallup survey, 57 percent of AI-using teachers said AI improves the quality of their grading and feedback, and we'd bet those are overwhelmingly the teachers who verify. The tool improves the work of people who check it. It embarrasses the people who don't.

Segment 5 · The lab: build for a real assessment (slides 18–24)

SLIDE 18 Lab setup

0:30

MS. RIVERA'S SCREEN · SO FAR Her lab: the assessment on her calendar, materials built tonight-free.

→ *Dark slide. Devices out, pairs formed inside two minutes. Announce the tool. This is the longest lab so far; protect every minute of it.*

Lab time, nineteen minutes, and today you build for real. Pick an assessment that's actually coming: a unit test, an essay, a project, a shop practical, anything in the next month or so. Grab your notes on what good work looks like, even if they're three messy lines. Ground rules: no student information of any kind; rubric first, then questions, then feedback starters; and before the lab ends, you catch and cut at least one weak question yourself. That last one is a requirement, not a joke. Go.

→ *If you grade with a rubric of your own, hold it up or put it on the screen here. One real rubric from a real room tells the lab what "done" looks like better than any slide.*

SLIDE 19 Lab step 1: the rubric

0:32

MS. RIVERA'S SCREEN · SO FAR On her screen: notes in, rubric out, hand-it-to-a-student pass next.

→ *6 minutes. Circulate. Weak rubrics are usually missing the teacher's notes; nudge with "what do you actually look for when you grade this? Type that."*

Step one, six minutes: notes to rubric. Use the workshop prompt from your handout, swap in your grade, subject, assignment, and criteria, paste your notes, run it. Then iterate with Kit 2's three rounds. Content round: are these the right criteria, and is anything you actually grade for missing? Voice round: the hand-it-to-a-student test, on every descriptor. Details round: table format, level names, point values if you use them. 45-min cut: run the voice round only; it's the one that matters most.

SLIDE 20 Lab step 2: the question bank

0:38

MS. RIVERA'S SCREEN · SO FAR Her bank: ten questions, misconception distractors, key checked.

→ 5 minutes. Listen for the misconception sentences; they're the best writing that happens all day.

Step two, five minutes: the bank. Same assessment, ten questions, mixed format, full answer key, distractor rationale. The one thing only you can supply: the misconception. Finish this sentence before you prompt: "my students keep confusing X with Y." Put that sentence in the prompt and make it a distractor. 45-min cut: six questions instead of ten.

SLIDE 21 Lab step 2, continued: the vetting rep

0:42

MS. RIVERA'S SCREEN · SO FAR Her vetting rep: found the weak question, caught and cut it.

Now the rep, two minutes, everyone: read your bank like a hostile editor. Check the answer key against your own knowledge. Check every distractor for accidental truth. Check the reading level. Find the weakest question, and either fix it by hand or delete it and write its replacement yourself. When you find it, and you will, say "got one" out loud. That little ritual is the whole safety system for this workflow: drafts are cheap, vetting is sacred.

SLIDE 22 Lab step 3: feedback starters

0:44

MS. RIVERA'S SCREEN · SO FAR Her starters: three patterns she writes every stack, drafted once.

→ 3–4 minutes. If someone starts pasting a student's work, redirect gently: describe the pattern instead; excerpts only if fully anonymous.

Step three: feedback starters. Name your three most common student patterns for this kind of work; you know them by heart. Run the pattern prompt: strength, pattern, next step, in plain student language. Then save all three where you grade, and paste your best one into the staff prompt doc's new Assessment section, with your name. These are starters, not stamps: in real grading you'll personalize every single one before it reaches a kid.

SLIDE 23 Share-out

0:46

MS. RIVERA'S SCREEN · SO FAR Her share-out: the catch story goes first.

→ 3–4 voices, one minute each. The weak-question stories go first; celebrate every catch. 45-min cut: two voices.

Let's hear it. First: who caught a weak question, and what was wrong with it? Those stories go first because that's the habit we're building. Then: read us one rubric descriptor that passes the hand-it-to-a-student test, and one feedback starter you'll actually use this week.

SLIDE 24 What you just built

0:48

MS. RIVERA'S SCREEN · SO FAR Her inventory: rubric, vetted bank, three starters, one caught question.

Look at the inventory from nineteen minutes: a four-level rubric a student can read, a vetted question bank with an answer key that's been checked by a professional, three feedback starters for the comments you write most, and one weak question caught and cut as proof the vetting reflex works. That set used to be a weekend. The judgment in it is all yours; only the typing got faster.

Segment 6 · Making it stick (slides 25–27)

SLIDE 25 Into the staff doc

0:49

MS. RIVERA'S SCREEN · SO FAR Her templates: filed in the staff doc's Assessment section.

Making it stick, part one: the staff prompt doc gets an Assessment section today, and your templates live there. The rubric prompt, the question-bank prompt, the starter prompt, each with blanks where the specifics go. The English department's rubric template will travel to science; the shop's safety-quiz bank prompt will travel to health class. Same commitment as Kit 2: when a prompt works, it goes in the doc. Nobody at this school solves the same problem twice.

SLIDE 26 The five hours

0:51

MS. RIVERA'S SCREEN · SO FAR Her five hours: the drafting share just moved to a faster tool.

Part two, the vision, stated carefully. The median teacher spends five hours a week on grading and feedback. We are not promising you those five hours back; grading still means reading real work and making real calls, and it should. But a real share of those hours was never judgment. It was drafting: the rubric written at 9 p.m., the quiz typed from scratch, the same comment keyed forty times. That share now moves to a machine that drafts and a teacher who decides. Whatever it gives you back, spend it on the part of assessment only you can do, or spend it on your own life. Both are wins.

SLIDE 27 What's next

0:52

MS. RIVERA'S SCREEN · SO FAR Her next kit: the student side. Integrity and clear expectations.

Next in the series: Kit 5, Academic Integrity. Today was the teacher side of assessment; Kit 5 is the student side: what honesty and original work look like when every student has these tools in their pocket too. Your rubric from today comes along; it turns out clear success criteria are one of the best integrity tools we have.

Segment 7 · First 48 hours & close (slides 28–30)

SLIDE 28 Your first 48 hours

0:54

MS. RIVERA'S SCREEN · SO FAR Her 48 hours: rubric out with the assignment, finish vetting, use one starter.

Three small things within two days, each under fifteen minutes: put the lab rubric in front of the students it was built for, when you give the assignment. Finish the vetting pass on your question bank, every question against the key, and file it where you'll find it. And the next time you grade a stack, use one of your three starters on a common pattern, personalized, and notice what it does to your pace. The sheet's at the door.

SLIDE 29 Exit ticket

0:55

MS. RIVERA'S SCREEN · SO FAR Her exit ticket: what she built, the weak question she caught.

→ *Distribute exit tickets; collect at the door. They're the school's PD documentation.*

Two minutes on the exit ticket: what you built, what you still want help with, and the weak question you caught. Yes, we're asking about the catch. It's the answer we're proudest of.

SLIDE 30 Close

0:56

Last word. Everything we built today is paper. The rubric is paper. The questions are paper. Good paper, faster paper, but paper. Knowing what a ten-year-old meant to say, and what they need to hear next: that was never paper. That's you. It stays you. Thanks, everyone.

© Adam & Katelyn Spinozzi · AI-Ready School. Licensed for use within the purchasing school. This resource is provided for informational purposes by educators, not attorneys, and is not legal advice. Consult your district's legal counsel before adopting any policy. Third-party tools named in this kit (ChatGPT, Claude, Gemini, Copilot) are referenced factually; no partnership or endorsement is implied.