

WHAT AI ACTUALLY COSTS

Somebody in your building has asked what it would cost to use AI with students, and the answer they got was a shrug. **Here is why nobody will give you a number, and how to work out your own** — a real one, for your class, in about twenty minutes.

2

completely different things
people mean by "cost"

1

worksheet that gets you a real
figure

0

prices printed here for you to
trust

The sentence this whole page is built on.

That is a measurement, not an estimate. Everyone else guesses. You can go and find out, and the finding out is easier than the guessing.

Why there are no prices in here

AI prices have moved several times a year, always downward, and a number printed in a PDF is wrong within months. **So this page teaches you where to look and what to multiply**, which does not go stale. Anybody who hands you a per-student figure without asking what your students will actually do with it is guessing.

Where this comes from

I built a button into the software I run for my own high school woodshop course that pushes one student's worth of work through for real and reports what it used. Not a projection — the actual figure. **This is that method, with the software taken out**, so you can do it with a calculator.

Two different questions wearing one word

Almost every useless answer about AI cost comes from mixing these up. **Work out which one you are being asked before you answer it.**

1 • A seat – what one person pays to use a chatbot

A monthly subscription per person. A free tier, then a paid tier at a fixed price a month, sometimes an education or team rate. **Predictable, and it scales by how many people, not by how much they use it.**

This is what most teachers are actually paying for today, and for a teacher-does-the-prompting classroom it is usually the whole bill.

2 • Usage – what a piece of software pays to call AI

Billed by the amount of text, per use, with no subscription. **This is what you pay if a tool you or your district built calls AI on its own** — grading drafts, generating practice, answering questions inside an app.

Unpredictable until you measure it, which is the entire reason for this page.

The question to ask back, before you promise anybody a number.

"Is a person typing into a chatbot, or is software calling AI by itself?" **Those have different bills, different risks and different permission questions**, and an answer to one is not an answer to the other.

And the cost nobody puts in the spreadsheet

Teacher time. An hour of checking what AI produced is a real cost and it is larger than the subscription. It is also the one that falls if the thing is set up well, and rises forever if it is not. Put it on the sheet, even as a guess, or the comparison is dishonest.

Five mechanics that do not go out of date

You do not need to know prices to reason about this. You need to know what gets counted.

1 • Text is counted in tokens, not words

A token is a chunk of text, usually a bit less than a word. Roughly, a page of writing is around 500 to 700 tokens. You do not have to count them yourself — the tools report the number.

2 • What goes in and what comes out are priced separately — and out costs much more

Across the major providers, generated text is billed at several times the rate of text you send in — commonly around five times. So asking for a long answer costs far more than sending a long document, which is the opposite of what most people assume.

3 • A conversation re-sends itself every turn

Each message in a back-and-forth sends the whole history again. A twenty-turn conversation is not twenty small requests; it is twenty requests that keep growing. This is the single biggest surprise on a first bill.

4 • Repeated context can be cached, and the discount is large

If the same long text is sent over and over — a rubric, a reading, a set of instructions — providers let software mark it as cached. Re-reading it then costs roughly a tenth of the normal input rate, at the price of a small surcharge the first time. For anything that sends the same rubric thirty times, this is the difference between affordable and not.

5 • Work that can wait is about half price

Providers offer a batch mode for work with no hurry — overnight rather than instantly — at roughly half the normal rate. Most classroom work that is not happening live qualifies, and almost nobody uses it.

Do not trust a price printed in any PDF, including this one.

There are none here on purpose. Look up the current rate on the provider's own pricing page on the day you do the arithmetic. The five mechanics above are what stay true.

Measure one, then multiply

Twenty minutes, and you end with a number you can defend in a meeting because you watched it happen.

1 Do the real thing once, for one student

Not a sample, not a shortened version. **One student's actual work, all the way through**, in whatever tool would really do it. An estimate built on a simplified version measures the simplified version.

2 Write down what it reported

Tokens in and tokens out, separately, because they are priced separately. Every serious tool reports both — in its usage screen, its billing page, or the response itself. **If a tool will not tell you, that is worth knowing too.**

3 Look up today's rate and do two multiplications

Tokens in times the input rate, tokens out times the output rate. Rates are quoted per million tokens, so **divide by a million at the end**. That is one student, once.

4 Multiply by reality, not by your class list

Students times **how many times each** — and the second number is the one people forget. A tool used once a term and a tool used every lesson differ by a factor of thirty, not by a rounding error.

5 Then ask the three questions that move the number

Is the same long text sent every time, so it could be cached? Does this need to be instant, or could it run overnight at half price? **Are we asking for long answers where a short one would do?** Output is the expensive half.

Why one real run beats any estimate.

The number depends almost entirely on how much text moves, and nobody can guess that about your assignment, your rubric and your students. **Twenty minutes of measuring beats an afternoon of arguing about a guess**, and it is the same twenty minutes whoever asks.

YOUR TURN

Work out your number

Fill this in while you do it, not afterwards. Keep it — when somebody asks next term, you will have the working, not just the answer.

What one student actually did

The task, in one line:

Tokens in

Tokens out

Today's rates, from the provider's own page

Input, per million

Output, per million

Date I looked

The arithmetic

One student, once

× students × times each

Per term

What would move it, and by how much

Cacheable repeated text? Could it run overnight? Are we asking for longer answers than we need?

If this is a conversation your whole staff needs to have

Our **AI Foundations & Safety** staff training kit is the whole-staff version — deck with speaker notes, a word-for-word script, a participant handout and an administrator one-pager, presentable cold by somebody who has never run a session. Licenses are half price. Search our store for it.

If this was useful, would you leave a review?

TPT gives you credit toward future purchases for every review, and it is the only thing that decides whether another teacher ever sees this.