



FLIP THROUGH BEFORE YOU DOWNLOAD

What AI Actually Costs

Three real pages, straight off the shipping files,
plus everything the full download includes.



Two different questions wearing one word

Almost every useless answer about AI cost comes from mixing these up. **Work out which one you are being asked before you answer it.**

1 • A seat — what one person pays to use a chatbot

A monthly subscription per person. A free tier, then a paid tier at a fixed price a month, sometimes an education or team rate. **Predictable, and it scales by how many people, not by how much they use it.**

This is what most teachers are actually paying for today, and for a teacher-does-the-prompting classroom it is usually the whole bill.

2 • Usage — what a piece of software pays to call AI

Billed by the amount of text, per use, with no subscription. **This is what you pay if a tool you or your district built calls AI on its own** — grading drafts, generating practice, answering questions inside an app.

Unpredictable until you measure it, which is the entire reason for this page.

The question to ask back before you can even pay a number.

"Is a person typing into a chatbot, or is software calling AI by itself?" **Those have different bills, different risks and different permissions for questions,** and an answer to one is not an answer to the other.

And the cost nobody puts in the spreadsheet

Teacher time. An hour of checking what AI produced is a real cost and it is larger than the subscription. It is also the one that falls if the thing is set up well, and rises forever if it is not. Put it on the sheet, even as a guess, or the comparison is dishonest.

SNEAK PEEK • 5 pages

Five mechanics that do not go out of date

You do not need to know prices to reason about this. You need to know what gets counted.

1 • Text is counted in tokens, not words

A token is a chunk of text, usually a bit less than a word. Roughly, a page of writing is around 500 to 700 tokens. You do not have to count them yourself — the tools report the number.

2 • What goes in and what comes out are priced separately — and out costs much more

Across the major providers, generated text is billed at several times the rate of text you send in — commonly around five times. So asking for a long answer costs far more than sending a long document, which is the opposite of what most people assume.

3 • A conversation re-sends itself every turn

Each message in a back-and-forth sends the whole history again. A twenty-turn conversation is not twenty small requests; it is twenty requests that keep growing. This is the single biggest surprise on a first bill.

4 • Repeated context can be cached, and the discount is large

If the same long text is sent over and over — a rubric, a reading, a set of instructions — providers let software mark it as cached. Re-reading it then costs roughly a tenth of the normal input rate, at the price of a small surcharge the first time. For anything that sends the same rubric thirty times, this is the difference between affordable and not.

5 • Work that can wait is about half price

Providers offer a batch mode for work with no hurry — overnight rather than instantly — at roughly half the normal rate. Most classroom work that is not happening live qualifies, and almost nobody uses it.

Do not trust a price printed in any PDF, including this one.

There are none here on purpose. Look up the current rate on the provider's own pricing page on the day you do the arithmetic. The five mechanics above are what stay true.

SNEAK PEEK • 5 pages

Measure one, then multiply

Twenty minutes, and you end with a number you can defend in a meeting because you watched it happen.

1 Do the real thing once, for one student

Not a sample, not a shortened version. **One student's actual work, all the way through**, in whatever tool would really do it. An estimate built on a simplified version measures the simplified version.

2 Write down what it reported

Tokens in and tokens out, separately, because they are priced separately. Every serious tool reports both — in its usage screen, its billing page, or the response itself. **If a tool will not tell you, that is worth knowing too.**

3 Look up today's rate and do two multiplications

Tokens in times the input rate, tokens out times the output rate. Rates are quoted per million tokens, so **divide by a million at the end**. That is one student, once.

4 Multiply by reality, not by your class list

Students times **how many times each** — and the second number is the one you forget. A tool used once a term and a tool used every lesson differ by a factor of thirty, not by a factor of three.

5 Then ask the three questions that matter

Is the same long text sent every time, or could be cached? Does this need to be instant, or could it run overnight at half price? Are we asking for long answers where a short one would do? Output is the expensive half.

SNEAK PEEK · 5 pages

Why one real run beats any estimate.

The number depends almost entirely on how much text moves, and nobody can guess that about your assignment, your rubric and your students. **Twenty minutes of measuring beats an afternoon of arguing about a guess**, and it is the same twenty minutes whoever asks.



WHAT IS INSIDE

What AI Actually Costs

- ✓ **Two different questions** wearing one word
- ✓ Five mechanics that do not go out of date
- ✓ **The five-step method** that gets you a real number
- ✓ The three questions that actually move the total
- ✓ **A fill-in worksheet** you can defend in a meeting

Running it with your whole staff? Additional teacher licenses are just \$0, and we accept purchase orders through TpT for Schools.

